

<https://towardsdatascience.com/sound-classification-using-images-68d4770df426>

▼ solution to ERRORS

1/10



Display the audio file

```
import IPython
IPython.display.Audio("/content/drive/My Drive/1.wav")
```

0:00 / 0:05

another good reading resource

Audio Classification using FastAI and On-the-Fly Frequency Transforms

<https://towardsdatascience.com/audio-classification-using-fastai-and-on-the-fly-frequency-transforms-4dbe1b540f89>

We have 2 options to convert the audio files to spectrograms, matplotlib or librosa. We will go for the latter because it is easier to use and well known in the sound domain. Before we use it we just need to install a little dependency to ensure librosa works well

I installed librosa because I wanted to use it but it didn't work so I used wavfile later.

```
pip install librosa
```

Double-click (or enter) to edit

The below code is for single channel or mono audio so if I use it for gravel classification it gives the following errors **ValueError: *only 1-dimensional arrays can be us*


<https://dsp.stackexchange.com/questions/10743/generating-spectrograms-in-python-with-less-noise>

```
import pylab
from scipy.io import wavfile

fs, frames = wavfile.read("/content/drive/My Drive/gravel.wav")

channels = [
    np.array(frames[:, 0]),
    np.array(frames[:, 1])
]
```

```
# generate specgram
Pxx, freqs, t, plot = pylab.specgram(
    channels[0],
    NFFT=4096,
    Fs=44100,
    detrend=pylab.detrend_none,
    window=pylab.window_hanning,
    noverlap=int(4096 * 0.5))
```




```
import pylab
from scipy.io import wavfile
```

```
fs, frames = wavfile.read("/content/drive/My Drive/nongravel.wav")
```

```
channels = [
    np.array(frames[:, 0]),
    np.array(frames[:, 1])
]
```

```
# generate specgram
Pxx, freqs, t, plot = pylab.specgram(
    channels[0],
    NFFT=4096,
    Fs=44100,
    detrend=pylab.detrend_none,
    window=pylab.window_hanning,
    noverlap=int(4096 * 0.5))
```



```
import wave, os, glob
import pylab
from scipy.io import wavfile
import matplotlib.pyplot as plt
import matplotlib

zero = []
path = '/content/drive/My Drive/audiofiles/gravel'
files = os.listdir(path)
gravel_counter= 0

for filename in glob.glob(os.path.join(path, '*.wav')):
    s = wave.open(filename, 'r')
    fs, frames = wavfile.read(filename)
    channels = [
        np.array(frames[:, 0]),
        np.array(frames[:, 1])
    ]

# generate specgram

    w = 10
    h = 7
    d = 70
    fig = plt.figure(figsize=(w, h), dpi=d)
    Pxx, freqs, t, plot = pylab.specgram(
        channels[0],
        NFFT=4096,
        Fs=44100,
        detrend=pylab.detrend_none,
        window=pylab.window_hanning,
        noverlap=int(4096 * 0.5))
    # print(filename)
#print(files)
    output_dir = "/content/drive/My Drive/Spectrogram/Gravel"
#fig.savefig('{}graph.png'.format(output_dir)) # old correct one
#print(s)

    gravel_counter = gravel_counter + 1
#print(gravel)
    gravel = "gravel" + str(gravel_counter) + ".png"
#print(filename)
    filepath = os.path.join(output_dir, gravel)

    plt.savefig('/content/drive/My Drive/Spectrogram/Gravel/' + gravel)
    plt.close()
```

```

    #print(filepath)
    #fig.savefig({}/filenamesave.format(output_dir))
#fig.savefig('{}graph'+'.png'.format(output_dir))
print("end")

    end

```

For Non Gravel sounds

```

import wave, os, glob
import pylab
from scipy.io import wavfile
import matplotlib.pyplot as plt
import matplotlib

zero = []
path = '/content/drive/My Drive/audiofiles/non_gravel'
files = os.listdir(path)
ngravel_counter= 0

for filename in glob.glob(os.path.join(path, '*.wav')):
    s = wave.open(filename, 'r')
    fs, frames = wavfile.read(filename)
    channels = [
        np.array(frames[:, 0]),
        np.array(frames[:, 1])
    ]

# generate specgram

    w = 10
    h = 7
    d = 70
    fig = plt.figure(figsize=(w, h), dpi=d)
    Pxx, freqs, t, plot = pylab.specgram(
        channels[0],
        NFFT=4096,
        Fs=44100,
        detrend=pylab.detrend_none,
        window=pylab.window_hanning,
        noverlap=int(4096 * 0.5))

    output_dir = "/content/drive/My Drive/Spectrogram/Non-gravel"

    ngravel_counter = ngravel_counter + 1

    ngravel = "Non-gravel" + str(ngravel_counter) + ".png"

    filepath = os.path.join(output_dir, gravel)

    plt.savefig('/content/drive/My Drive/Spectrogram/Non-gravel/' + ngravel)

```

```
plt.close()
```

```
print("end")
```

```
end
```

classification part


```
path = "/content/drive/My Drive/Spectrogram";
```

```
path
```


```
'/content/drive/My Drive/Spectrogram'
```

```
data = (ImageList.from_folder(path)
        .split_by_rand_pct()
        .label_from_folder()
        .transform([], size = 224)
        .databunch())
data.show_batch(rows=3, figsize=(10,10))
```


Gravel



Gravel



Gravel



```
data.normalize(imagenet_stats)
```

Gravel

Non-gravel

Non-gravel

```
learn = cnn_learner(data, models.resnet34, metrics=[error_rate, accuracy])
```

Downloading: "<https://download.pytorch.org/models/resnet34-333f7ec4.pth>" to /root/.cache/torch/hub/ - 100%
83.3M/83.3M [04:38<00:00, 313kB/s]



```
learn.model
```

```
learn.fit_one_cycle(10)
```

epoch	train_loss	valid_loss	error_rate	accuracy	time
0	1.292108	1.744478	0.340426	0.659574	00:16
1	1.195238	1.300250	0.340426	0.659574	00:07
2	0.927550	0.491724	0.212766	0.787234	00:06
3	0.735525	0.387641	0.170213	0.829787	00:06
4	0.607243	0.362773	0.234043	0.765957	00:06
5	0.501051	0.496762	0.191489	0.808511	00:06
6	0.422047	0.555415	0.170213	0.829787	00:06
7	0.369687	0.533067	0.148936	0.851064	00:06
8	0.323575	0.504000	0.106383	0.893617	00:06
9	0.286102	0.406760	0.127660	0.872340	00:06

```
learn.save('stage-1')
```

```
interp = ClassificationInterpretation.from_learner(learn)
```

```
losses,idxs = interp.top_losses()
```


```
len(data.valid_ds)==len(losses)==len(idxs)
```

```
True
```


```
interp.plot_top_losses(9, figsize=(15,11))
```

Prediction/Actual/Loss/Probability


Gravel/Non-gravel / 6.42 / 0.00




Non-gravel/Gravel / 4.29 / 0.01




Gravel




Gravel/Non-gravel / 2.33 / 0.10




Gravel/Non-gravel / 1.52 / 0.22




Non-gr




Non-gravel/Non-gravel / 0.51 / 0.60



Non-gravel/Non-gravel / 0.24 / 0.79



Grav



```
interp.plot_confusion_matrix(figsize=(12,12), dpi=60)
```


```
learn.unfreeze()
```

```
learn.lr_find()
```


48.00% [24/50 02:02<02:12]

epoch	train_loss	valid_loss	error_rate	accuracy	time
0	0.030606	#na#	00:05		
1	0.018546	#na#	00:05		
2	0.016112	#na#	00:05		
3	0.018568	#na#	00:04		
4	0.017463	#na#	00:05		
5	0.018049	#na#	00:05		
6	0.016334	#na#	00:05		
7	0.017323	#na#	00:04		
8	0.016446	#na#	00:04		
9	0.015740	#na#	00:05		
10	0.014294	#na#	00:05		
11	0.013444	#na#	00:05		
12	0.012502	#na#	00:04		
13	0.011407	#na#	00:04		
14	0.010519	#na#	00:05		
15	0.010183	#na#	00:05		
16	0.009732	#na#	00:05		
17	0.009065	#na#	00:05		
18	0.008474	#na#	00:05		
19	0.007895	#na#	00:05		
20	0.007943	#na#	00:05		
21	0.008245	#na#	00:05		
22	0.015109	#na#	00:05		
23	0.018493	#na#	00:05		

```
learn.recorder.plot()
```



```
learn.unfreeze()  
learn.fit_one_cycle(10, max_lr=slice(1e-6,1e-5))
```

epoch	train_loss	valid_loss	error_rate	accuracy	time
0	0.027285	0.299077	0.085106	0.914894	00:07
1	0.036268	0.245371	0.127660	0.872340	00:07
2	0.028660	0.192131	0.063830	0.936170	00:07
3	0.027103	0.157950	0.063830	0.936170	00:07
4	0.022239	0.141795	0.042553	0.957447	00:07
5	0.020072	0.168879	0.042553	0.957447	00:07
6	0.017383	0.215570	0.063830	0.936170	00:07
7	0.015289	0.262880	0.063830	0.936170	00:07
8	0.013922	0.307496	0.063830	0.936170	00:07
9	0.013757	0.336888	0.063830	0.936170	00:07